

kafka-isotope Library

kafka-isotope provides **end-to-end record tracing for Apache Kafka** using Kafka record headers, a **ProducerInterceptor**, and optional **Prometheus metrics** and **OpenTelemetry (OTel) spans**.

It works by attaching lightweight tracing artifacts—called *isotopes*—to records as they move through Kafka event pipelines.

An isotope is a small tracing payload carried in Kafka record headers. Like a biochemical isotope used to trace molecules through a metabolic pathway, it enables the journey of a record through an event-driven architecture to be observed and analyzed.

On first produce, **kafka-isotope** stamps a UUIDv7 trace ID and origin metadata onto the record. Each subsequent re-produce appends another hop, allowing the full path of a record through Kafka pipelines to be reconstructed.

Built entirely on Kafka's public extension points—a **ProducerInterceptor** plus record headers—**kafka-isotope** requires:

- No broker changes
- No sidecar agents
- No vendor lock-in

From isotope headers, seven reports can be derived:

- End-to-end latency
- Latency percentiles (p50/p95/p99)
- Pipeline topology
- Bipartite topology (producer → topic → consumer)
- Hop distribution
- Coverage (where traces drop off)
- Stuck trace detection

Three of them (end-to-end latency, pipeline topology and hop distribution) can be exported directly as **Prometheus metrics** with the **kafka-isotope-metrics** module. All seven run as Flink SQL in the companion **confluent-kafka-isotope** demo.

Separately, the optional **kafka-isotope-otel** module emits each hop as an **OTel span**, so the same traces can be viewed in any **OpenTelemetry Protocol (OTLP)** backend alongside the rest of your distributed tracing.

This makes **kafka-isotope** a lightweight but powerful observability layer for Kafka-based event-driven architectures.

Table of Contents


- **1.0 Modules**
- **2.0 Install using Gradle**

- [3.0 Demo](#)

1.0 Modules

Module	Coordinate	Pulls	Use it for
kafka-isotope-core	<code>ai.signalroom:kafka-isotope-core</code>	Jackson, SLF4J (kafka-clients is <code>compileOnly</code>)	Trace propagation — interceptor, headers, consume markers.
kafka-isotope-metrics	<code>ai.signalroom:kafka-isotope-metrics</code>	kafka-isotope-core + Micrometer/Prometheus	Optional <code>/metrics</code> exporter for the stateless reports.
kafka-isotope-otel	<code>ai.signalroom:kafka-isotope-otel</code>	kafka-isotope-core + OTLP protobuf	Optional span writer — one OTel span per hop on a Kafka topic, read by a stock OTel Collector.

[kafka-isotope-core](#) has no metrics or tracing dependency. It routes every emission through two sinks that start out as no-ops (`NoOpMetricsSink` and `NoOpSpanSink`). Adding [kafka-isotope-metrics](#) or [kafka-isotope-otel](#) to the classpath doesn't change that: the real sink is registered only when your application calls `PrometheusIsotopeMetrics.start(port)` or `KafkaOtlpSpanSink.start(config)`, and stays a no-op if that call fails. Until then, each record costs an "is it enabled?" check and no metric or span work.



Dashed arrows are optional: the Prometheus metrics and OTel span modules only run after their **start** call, and Flink SQL lives in the companion **confluent-kafka-isotope** demo repository, not in this library.

2.0 Install using Gradle

```
repositories {
    mavenCentral()
}

dependencies {
    implementation 'ai.signalroom:kafka-isotope-core:0.19.0'
    implementation 'ai.signalroom:kafka-isotope-metrics:0.19.0' // optional
    - only for Prometheus metrics
    implementation 'ai.signalroom:kafka-isotope-otel:0.19.0' // optional
    - only for OTel spans
}
```

3.0 Demo

To showcase the capabilities of the `kafka-isotope` library, the companion `confluent-kafka-isotope` repository provides a runnable reference e-commerce order event pipeline that can be deployed on Confluent Platform on minikube or Confluent Cloud. Both runtimes support all seven Flink reports and an optional Prometheus + Grafana stack (hosted on minikube) for metrics collection and visualization.